

CS543 – Assignment 4

Pallaw Kumar (pallawk2)

Question 1.

1.1

model by using Python print command:

```
BaseNet(  
    (conv1): Conv2d(1, 6, kernel_size=(5, 5), stride=(1, 1))  
    (relu1): ReLU()  
    (pool1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,  
ceil_mode=False)  
    (conv2): Conv2d(6, 12, kernel_size=(5, 5), stride=(1, 1))  
    (relu2): ReLU()  
    (pool2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,  
ceil_mode=False)  
    (fc1): Linear(in_features=192, out_features=12, bias=True)  
    (relu3): ReLU()  
    (fc2): Linear(in_features=12, out_features=10, bias=True)  
)
```

Final accuracy on the validation set:

```
Accuracy of the final network on the val images: 87.0 %  
Accuracy of T-shirt/top : 90.6 %  
Accuracy of Trouser : 98.4 %  
Accuracy of Pullover : 71.8 %  
Accuracy of Dress : 83.3 %  
Accuracy of Coat : 80.6 %  
Accuracy of Sandal : 97.6 %  
Accuracy of Shirt : 61.0 %  
Accuracy of Sneaker : 93.3 %  
Accuracy of Bag : 96.1 %  
Accuracy of Ankle boot : 95.6 %
```

Question 1.2

Best model:

```
BaseNet(  
    (conv1): Conv2d(1, 32, kernel_size=(5, 5), stride=(1, 1))  
    (bn1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,  
track_running_stats=True)  
    (conv2): Conv2d(32, 64, kernel_size=(5, 5), stride=(1, 1))  
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,  
track_running_stats=True)  
    (fc1): Linear(in_features=1024, out_features=512, bias=True)  
    (bn3): BatchNorm1d(512, eps=1e-05, momentum=0.1, affine=True,  
track_running_stats=True)  
    (fc2): Linear(in_features=512, out_features=256, bias=True)  
    (bn4): BatchNorm1d(256, eps=1e-05, momentum=0.1, affine=True,  
track_running_stats=True)  
    (fc3): Linear(in_features=256, out_features=128, bias=True)  
    (bn5): BatchNorm1d(128, eps=1e-05, momentum=0.1, affine=True,  
track_running_stats=True)  
    (fc5): Linear(in_features=128, out_features=10, bias=True)  
    (pool1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,  
ceil_mode=False)  
    (pool2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,  
ceil_mode=False)  
    (drop1): Dropout(p=0.2, inplace=False)  
    (drop2): Dropout(p=0.2, inplace=False)  
    (relu): ReLU()  
)
```

Final accuracy on validation set:

Accuracy of the final network on the val images: 92.1 %
Accuracy of T-shirt/top : 87.3 %
Accuracy of Trouser : 97.8 %
Accuracy of Pullover : 87.1 %
Accuracy of Dress : 93.3 %
Accuracy of Coat : 86.3 %
Accuracy of Sandal : 97.6 %
Accuracy of Shirt : 79.4 %
Accuracy of Sneaker : 96.9 %
Accuracy of Bag : 99.2 %
Accuracy of Ankle boot : 95.8 %


Table defining final architecture:

Layer No	Layer Type	Kernel size	Input Dim	Output Dim	Input Channels	Output Channels
1	Conv2d	5	28	24	1	32
2	BatchNorm2d	-	24	24	32	32
3	relu	-	24	24	32	32
4	Maxpool2d	2	24	12	32	32
5	Dropout	-	12	12	32	32
6	Conv2d	5	12	8	32	64
7	BatchNorm2d	-	8	8	64	64

8	relu	-	8	8	64	64
9	Maxpool2d	2	8	4	64	64
10	linear	-	1	1	1024	512
11	BatchNorm1d	-	1	1	1024	
12	relu	-	1	1	512	512
13	Dropout	-	1	1	512	512
14	linear	-	1	1	512	256
15	BatchNorm1d	-	1	1	256	256
16	relu	-	1	1	256	256
17	linear	-	1	1	256	128
18	BatchNorm1d	-	1	1	128	128
19	relu	-	1	1	128	128
20	linear	-	1	1	128	10

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 24, 24]	832
BatchNorm2d-2	[-1, 32, 24, 24]	64
ReLU-3	[-1, 32, 24, 24]	0
MaxPool2d-4	[-1, 32, 12, 12]	0
Dropout-5	[-1, 32, 12, 12]	0
Conv2d-6	[-1, 64, 8, 8]	51,264
BatchNorm2d-7	[-1, 64, 8, 8]	128
ReLU-8	[-1, 64, 8, 8]	0
MaxPool2d-9	[-1, 64, 4, 4]	0
Linear-10	[-1, 512]	524,800
BatchNorm1d-11	[-1, 512]	1,024
ReLU-12	[-1, 512]	0
Dropout-13	[-1, 512]	0
Linear-14	[-1, 256]	131,328
BatchNorm1d-15	[-1, 256]	512
ReLU-16	[-1, 256]	0
Linear-17	[-1, 128]	32,896
BatchNorm1d-18	[-1, 128]	256
ReLU-19	[-1, 128]	0
Linear-20	[-1, 10]	1,290
Total params: 744,394		
Trainable params: 744,394		
Non-trainable params: 0		
Input size (MB): 0.00		
Forward/backward pass size (MB): 0.62		
Params size (MB): 2.84		
Estimated Total Size (MB): 3.46		

Training loss plot and test accuracy plot for final model:



Ablation table:

Model Components	Validation Accuracy
Baseline Model	87.0%
+ Data Augmentation	87.7 %
+ Batch Normalization	89.0 %
+ Dropout	88.8 %
+ Increased Model Depth	89.7 %
+ Increased Model Width	92.1 %

Test set accuracy (category-wise and aggregate) for best model from gradescope:

Accuracy: 91.2 %
 Accuracy of T-shirt/top : 86.9 %
 Accuracy of Trouser : 98.4 %
 Accuracy of Pullover : 86.4 %
 Accuracy of Dress : 88.4 %
 Accuracy of Coat : 86.6 %
 Accuracy of Sandal : 96.9 %
 Accuracy of Shirt : 76.0 %
 Accuracy of Sneaker : 98.1 %
 Accuracy of Bag : 99.5 %
 Accuracy of Ankle boot : 97.4 %

Question 2:

2.1


To implement upsampling I have used `torch.nn.functional.interpolate` function for resizing a tensor to a specified size using interpolation. In this case, the function is used to resize the output tensor `x` to a shape of (224, 224) using bilinear interpolation. The mode parameter 'bilinear' is used which performs linear interpolation between the 4 nearest neighbors of each output pixel.

Layer No	Layer Type	Kernel size	Input Dim	Output Dim	Input Channels	Output Channels
1-17**	Pretrained ResNet18	-	224	7	3	512
19	Conv2d	-	7	7	512	6
20	Upsample	-	7	224	6	6

Training details:

- The training loop was changed such that the training loss, mean AP, and iou (for the validation set) could be displayed or tracked for each iteration/epoch.
- A $1e-4$ learning rate was applied. Forty epochs were used to train the model. At epoch 30, the best model was discovered.

visualization of validation set:



```
performance on validation set
background: AP: 0.93, IoU: 0.75
sports: AP: 0.51, IoU: 0.22
accessory: AP: 0.56, IoU: 0.33
animal: AP: 0.90, IoU: 0.67
vehicle: AP: 0.75, IoU: 0.50
person: AP: 0.83, IoU: 0.59
mean: AP: 0.75, IoU: 0.51
```

question 2.2

Best model:

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 112, 112]	9,408
BatchNorm2d-2	[-1, 64, 112, 112]	128
ReLU-3	[-1, 64, 112, 112]	0
MaxPool2d-4	[-1, 64, 56, 56]	0
Conv2d-5	[-1, 64, 56, 56]	36,864
BatchNorm2d-6	[-1, 64, 56, 56]	128
ReLU-7	[-1, 64, 56, 56]	0
Conv2d-8	[-1, 64, 56, 56]	36,864
BatchNorm2d-9	[-1, 64, 56, 56]	128
ReLU-10	[-1, 64, 56, 56]	0
BasicBlock-11	[-1, 64, 56, 56]	0
Conv2d-12	[-1, 64, 56, 56]	36,864
BatchNorm2d-13	[-1, 64, 56, 56]	128
ReLU-14	[-1, 64, 56, 56]	0
Conv2d-15	[-1, 64, 56, 56]	36,864
BatchNorm2d-16	[-1, 64, 56, 56]	128
ReLU-17	[-1, 64, 56, 56]	0
BasicBlock-18	[-1, 64, 56, 56]	0
Conv2d-19	[-1, 128, 28, 28]	73,728
BatchNorm2d-20	[-1, 128, 28, 28]	256
ReLU-21	[-1, 128, 28, 28]	0
Conv2d-22	[-1, 128, 28, 28]	147,456
BatchNorm2d-23	[-1, 128, 28, 28]	256
Conv2d-24	[-1, 128, 28, 28]	8,192
BatchNorm2d-25	[-1, 128, 28, 28]	256
ReLU-26	[-1, 128, 28, 28]	0
BasicBlock-27	[-1, 128, 28, 28]	0
Conv2d-28	[-1, 128, 28, 28]	147,456
BatchNorm2d-29	[-1, 128, 28, 28]	256
ReLU-30	[-1, 128, 28, 28]	0
Conv2d-31	[-1, 128, 28, 28]	147,456
BatchNorm2d-32	[-1, 128, 28, 28]	256

ReLU-33	[-1, 128, 28, 28]	0
BasicBlock-34	[-1, 128, 28, 28]	0
Conv2d-35	[-1, 256, 14, 14]	294,912
BatchNorm2d-36	[-1, 256, 14, 14]	512
ReLU-37	[-1, 256, 14, 14]	0
Conv2d-38	[-1, 256, 14, 14]	589,824
BatchNorm2d-39	[-1, 256, 14, 14]	512
Conv2d-40	[-1, 256, 14, 14]	32,768
BatchNorm2d-41	[-1, 256, 14, 14]	512
ReLU-42	[-1, 256, 14, 14]	0
BasicBlock-43	[-1, 256, 14, 14]	0
Conv2d-44	[-1, 256, 14, 14]	589,824
BatchNorm2d-45	[-1, 256, 14, 14]	512
ReLU-46	[-1, 256, 14, 14]	0
Conv2d-47	[-1, 256, 14, 14]	589,824
BatchNorm2d-48	[-1, 256, 14, 14]	512
ReLU-49	[-1, 256, 14, 14]	0
BasicBlock-50	[-1, 256, 14, 14]	0
Conv2d-51	[-1, 512, 7, 7]	1,179,648
BatchNorm2d-52	[-1, 512, 7, 7]	1,024
ReLU-53	[-1, 512, 7, 7]	0
Conv2d-54	[-1, 512, 7, 7]	2,359,296
BatchNorm2d-55	[-1, 512, 7, 7]	1,024
Conv2d-56	[-1, 512, 7, 7]	131,072
BatchNorm2d-57	[-1, 512, 7, 7]	1,024
ReLU-58	[-1, 512, 7, 7]	0
BasicBlock-59	[-1, 512, 7, 7]	0
Conv2d-60	[-1, 512, 7, 7]	2,359,296
BatchNorm2d-61	[-1, 512, 7, 7]	1,024
ReLU-62	[-1, 512, 7, 7]	0
Conv2d-63	[-1, 512, 7, 7]	2,359,296
BatchNorm2d-64	[-1, 512, 7, 7]	1,024
ReLU-65	[-1, 512, 7, 7]	0
BasicBlock-66	[-1, 512, 7, 7]	0
Conv2d-67	[-1, 512, 7, 7]	2,359,808
BatchNorm2d-68	[-1, 512, 7, 7]	1,024

```

ReLU-69          [-1, 512, 7, 7]          0
Upsample-70      [-1, 512, 14, 14]         0
Conv2d-71        [-1, 512, 14, 14]       2,359,808
BatchNorm2d-72   [-1, 512, 14, 14]         1,024
ReLU-73          [-1, 512, 14, 14]         0
Upsample-74      [-1, 512, 56, 56]         0
Conv2d-75        [-1, 512, 56, 56]       2,359,808
BatchNorm2d-76   [-1, 512, 56, 56]         1,024
ReLU-77          [-1, 512, 56, 56]         0
Conv2d-78        [-1, 6, 56, 56]          3,078
Upsample-79      [-1, 6, 224, 224]         0
=====
Total params: 18,262,086
Trainable params: 18,262,086
Non-trainable params: 0
-----
Input size (MB): 0.57
Forward/backward pass size (MB): 117.86
Params size (MB): 69.66
Estimated Total Size (MB): 188.10
-----

```

Final performance on validation set:

The final performance on the validation set was 0.79 mAP and 0.58 mIoU

```

performance on validation set
background: AP: 0.93, IoU: 0.75
  sports: AP: 0.61, IoU: 0.39
accessory: AP: 0.63, IoU: 0.41
  animal: AP: 0.93, IoU: 0.71
  vehicle: AP: 0.77, IoU: 0.60
  person: AP: 0.85, IoU: 0.62
    mean: AP: 0.79, IoU: 0.58

```

Model Components	Mean AP	mean IoU
Baseline Model	0.74	0.51
+ Data Augmentation	0.75	0.52
+ dilated convolution layers upsampling by 32	0.76	0.54
+ Batch Normalization	0.76	0.55
+ Increased Model Depth by adding more dilated convolution layers and upsampling step by step	0.79	0.58

Test set accuracy of best model from gradescope:

Q2 evaluation results:

mean: IoU: 0.53

mean: AP: 0.73

background: AP: 0.93, IoU: 0.73

sports: AP: 0.39, IoU: 0.28

accessory: AP: 0.63, IoU: 0.41

animal: AP: 0.83, IoU: 0.59

vehicle: AP: 0.75, IoU: 0.54

person: AP: 0.83, IoU: 0.60

Visualization:

